

ComicCraft: AI Comic Story Creator using Gemini Models

Date	30 September 2026
Team ID	[ENTER TEAM ID]
Project Name	ComicCraft - AI Comic Story Creator using Gemini Models
Maximum Marks	Project Documentation

Project Description

ComicCraft is a web-based application that uses AI to generate personalised comic stories and illustrations from user prompts. Built with FastAPI and integrated with Google's Gemini models and Stable Diffusion, it automates storylines, dialogue and comic-style imagery. Users give a story prompt, main character name, setting, tone and art style; ComicCraft returns a panel-by-panel story with illustrations, shown in the browser and downloadable as a PDF.

It offers both an HTML interface and JSON API endpoints, and uses Hugging Face Diffusers for image generation, so that every comic feels cohesive and visually engaging even for non-artists.

Scenarios

Scenario 1: A user enters "A brave fox exploring an enchanted forest", names the hero Rusty, and chooses forest, dramatic and anime. ComicCraft uses Gemini Flash for a 5-panel outline, Gemini Pro for narration and dialogue and Stable Diffusion for the artwork, producing a personalised comic.

Scenario 2: A user wants something light-hearted. On the form they select *funny* tone and *comic book* art style. The whole pipeline (outline, story and images) runs again with prompts tailored to a humorous narrative and a classic comic-book look, so users can iterate until the mood matches their vision.

Scenario 3: After reviewing the preview the user downloads the comic. `layout_builder.py` assembles each panel's title, image and text; `exporters.py` compiles them with FPDF into a timestamped PDF; the user then lands on the export success page.

Technical Architecture

ComicCraft has three parts: a **frontend** (HTML, CSS, Jinja2) that collects input and shows comics, a **backend** (FastAPI) that routes requests and orchestrates the workflow, and **AI integration** (Gemini via google-genai, Stable Diffusion via Diffusers). See 3. Project Design Phase/Solution Architecture.pdf for the diagram.

Pre-requisites

- FastAPI knowledge - <https://fastapi.tiangolo.com>
- Uvicorn ASGI server - <https://www.uvicorn.org>
- Python proficiency - <https://docs.python.org/3/>
- HTML, CSS and Jinja2 - <https://www.w3schools.com> and <https://jinja.palletsprojects.com>
- Google Gemini API - <https://ai.google.dev>
- Hugging Face Diffusers and Stable Diffusion - <https://huggingface.co/docs/diffusers>
- PyTorch basics - <https://pytorch.org/docs/stable/index.html>
- PDF generation with FPDF2 - <https://py-pdf.github.io/fpdf2/>

- Git - <https://git-scm.com/docs>
- Python environments (venv, pip) - <https://docs.python.org/3/library/venv.html>

Project Workflow

Milestone	Activities
1. Model selection and architecture	1.1 Research and select models; 1.2 Define architecture; 1.3 Set up environment
2. Core functionalities development	2.1 Develop core functions; 2.2 Implement FastAPI backend
3. routes.py development	3.1 Main application logic in routes.py
4. Frontend development	4.1 Design and build UI; 4.2 Dynamic Jinja2 templates
5. Deployment	5.1 Prepare for local deployment; 5.2 Test and verify

Milestone 1: Model Selection and Architecture

Activity 1.1: Research and Select the Appropriate Generative AI Model

ComicCraft must produce (a) a structured comic outline, (b) narration and dialogue and (c) an illustration per panel. We compared Gemini Flash (fast structured output), Gemini Pro (detailed creative text) and Stable Diffusion (image generation) on response quality, prompt adherence, speed, stability and documentation.

Role	Model chosen	Reason
Panel outline	Gemini Flash	Fast, follows JSON structure reliably, low cost
Narration and dialogue	Gemini Pro	Richer creative writing and consistent tone
Illustrations	Stable Diffusion v1.5 (Diffusers)	Open model, controllable styles, strong community support

Note on model names: the original brief lists models/gemini-1.5-flash, models/gemini-1.5-pro and runwayml/stable-diffusion-v1-5. Google has retired Gemini 1.5 and the SD 1.5 weights moved to stable-diffusion-v1-5/stable-diffusion-v1-5, so the code reads GEMINI_FLASH_MODEL, GEMINI_PRO_MODEL and SD_MODEL_ID from .env (defaults: gemini-flash-latest, gemini-pro-latest, stable-diffusion-v1-5/stable-diffusion-v1-5). The roles are the same as in the brief.

Activity 1.2: Define the Architecture of the Application

- **Frontend:** index.html (input), comic_preview.html (result), export_success.html (confirmation).
- **Backend:** FastAPI routes in routes.py: form and JSON generation, PDF download, export success, image testing.
- **AI integration:** Gemini Flash creates the 5-panel outline; Gemini Pro writes the story; Stable Diffusion draws the panels; results flow to templates and PDF export.

Activity 1.3: Set Up the Development Environment

```
python -m venv comiccraft-env
comiccraft-env\Scripts\activate           # Windows
pip install -r requirements.txt           # fastapi uvicorn jinja2 python-multipart google-genai diffusers transformers fpdf2
uvicorn app.main:app --reload             # http://127.0.0.1:8000 and /docs
```

Project structure:

```
app/ main.py routes.py config.py schemas.py services/{gemini_flash,gemini_pro,image_generator,layout_builder,exporter}
templates/ index.html comic_preview.html export_success.html
static/ img/ panels/ exports/ tests/ scripts/ requirements.txt .env.example
```

Milestone 2: Core Functionalities Development

Activity 2.1: Develop the Core Functionalities

Function	File	What it does
generate_outline()	services/gemini_flash.py	Asks Gemini Flash for a JSON array of 5 panels (panel_number, title, scene_description, image_prompt); cleans and validates the JSON.
generate_story()	services/gemini_pro.py	Asks Gemini Pro to turn the outline into caption and narration/dialogue for every panel; returns one formatted text block.
generate_image()	services/image_generator.py	Runs Stable Diffusion with the panel prompt plus art-style keywords, sanitises the file name, saves to static/panels and returns the path.
build_comic_layout()	services/layout_builder.py	Parses the story text and matches each panel with its image: list of dicts with panel number, image path and text.
save_pdf()	services/exporters.py	Builds a multi-page PDF with FPDF2 (cover + one page per panel) in static/exports with a timestamped name.

Example: outline generation (services/gemini_flash.py)

```
def generate_outline(story_prompt, character_name, setting, tone, art_style) -> list[dict]:
    prompt = OUTLINE_PROMPT.format(n=config.NUM_PANELS, story_prompt=story_prompt, ...)
    text = call_gemini(config.GEMINI_FLASH_MODEL, prompt, json_mode=True)
    return normalize_outline(_extract_json(text))    # always exactly 5 complete panels
```

Activity 2.2: Implement the FastAPI Backend

All routing is in routes.py. The HTML form is read with Form(...) parameters and the JSON API uses the PromptRequest Pydantic schema. Both call the same create_comic() function (outline, story, images, layout, PDF).

```
class PromptRequest(BaseModel):
    story_prompt: str = Field(..., min_length=5, max_length=500)
    character_name: str = Field("Hero", min_length=1, max_length=40)
    setting: str; tone: str; art_style: str
```

Milestone 3: routes.py Development

Activity 3.1: Writing the Main Application Logic in routes.py

Route	Method	Purpose
/	GET	Homepage with the story form (index.html)
/generate	POST	Form submission: run the AI workflow and render comic_preview.html
/generate-comic/json	POST	JSON API: same workflow, returns layout data and the PDF path
/download/{filename}	GET	Serve the exported PDF (file name checked against comic_*.pdf)
/export-success	GET	Confirmation page after the download
/test-image	GET	Developer utility: test image generation from a prompt
/health	GET	Liveness check

Error handling: invalid input returns HTTP 422 with a readable message; AI failures raise GenerationError and return 502; unexpected errors return 500 without exposing a stack trace. Handlers are plain def functions so FastAPI runs the slow AI work in a thread pool without blocking the server.

```
@router.post("/generate", response_class=HTMLResponse)
def generate(request: Request, story_prompt: str = Form(...), character_name: str = Form("Hero"), ...):
    req = PromptRequest(**form)    # validation
    layout, pdf_path = create_comic(req)    # outline -> story -> images -> layout -> PDF
    return templates.TemplateResponse(request, "comic_preview.html", {"layout": layout, "pdf_file": ..., "req": req})
```

Milestone 4: Frontend Development

Activity 4.1: Designing and Developing the User Interface

The interface uses HTML, embedded CSS and Jinja2. index.html has labelled fields for story prompt, character name, setting (school, forest, space, city), tone (light-hearted, dramatic, poetic, funny) and art style (anime, pixel art, comic book, realistic) on a scenic background. Buttons have hover effects, panels use rounded corners and shadows, and the layout is centred and responsive.

Activity 4.2: Creating Dynamic Templates with FastAPI's Jinja2

```
templates = Jinja2Templates(directory="templates")
{% for panel in layout %} <h2>Panel {{ panel.panel_number }}: {{ panel.title }}</h2>
   <p class="scene">{{ panel.scene_description }}</p> ... {% endfor %}
```

Form field names match the Form(...) parameters in FastAPI; comic_preview.html loops over the layout list; export_success.html receives the PDF file name.

Milestone 5: Deployment

Activity 5.1: Preparing the Application for Local Deployment

Create a virtual environment, install requirements.txt and store secrets in a .env file loaded with python-dotenv:

```
GEMINI_API_KEY=your_gemini_api_key_here
HF_API_KEY=your-huggingface-api-key-here
```

Activity 5.2: Testing and Verifying Local Deployment


```
uvicorn app.main:app --reload          # open http://127.0.0.1:8000 and http://127.0.0.1:8000/docs
```

Verification performed: 21 automated tests pass (services and routes), and four load-test scenarios ran without errors. These runs used the offline mock mode (COMICCRAFT MOCK=1) that replaces Gemini and Stable Diffusion with deterministic stand-ins; the screenshots below therefore show placeholder panels. Run the app with your own keys to see real AI output.

Exploring the Website's Web Pages


Home page

The starting point: a clean form on a scenic background where users enter the story prompt, main character name, setting, tone and art style.

The image shows a web form for 'ComicCraft' centered on a scenic background of stylized mountains, a sun, and a forest. The form is white with a red header 'ComicCraft' and a subtitle 'Turn a simple idea into a 5-panel AI comic - story, art and PDF.' Below this are several input fields: 'Story prompt' with the text 'A brave fox exploring an enchanted forest', 'Main character name' with 'Rusty', 'Setting' with a dropdown menu showing 'Forest', 'Story tone' with a dropdown menu showing 'Light-hearted', and 'Art style' with a dropdown menu showing 'Anime'. At the bottom of the form is a red button labeled 'Create my comic'.


Comic Preview (result) page

Each panel shows a numbered title, the illustration, the scene description in italics, the caption, the narration and the image prompt used. (Screenshot from offline mock mode.)




Download button and exported PDF

The "Download Your Comic as PDF" button downloads a PDF with a cover page and one page per panel (panel title, image, scene description, caption, narration).



Comic Export Success page

After the download starts, users see a confirmation and a "Go Create Another Comic" button.



Conclusion

ComicCraft shows how large language models (Gemini Flash and Pro) and Stable Diffusion can work together to turn a simple prompt into a complete illustrated comic with an on-screen preview and a downloadable PDF. Its modular design (one module per workflow step, a shared pipeline and configurable models) makes future features easy to add, such as user accounts, a comic library, character consistency, voice narration and cloud deployment.